

# YASH JAIN

Brooklyn, NY, 11220 | (917)-203-7016 | yj3273@stern.nyu.edu | linkedin.com/in/jainyash20 | github.com/yashjain20

## SUMMARY

Data Engineer with an MS in Information Systems from NYU and hands-on experience architecting production-grade data pipelines in cloud environments. Skilled in building scalable ETL/ELT workflows using **Python** and **SQL**, designing dimensional data models, and enforcing data quality at scale. Experienced with **AWS**, **Spark**, **Kafka**, and modern orchestration frameworks to deliver reliable, analytics-ready datasets.

## EDUCATION

### Master of Science, Information Systems

New York University, New York, NY

May 2026

GPA: 3.9/4.0

*Graduate Teaching Assistant:* Database Systems, Fundamental Algorithms

### Bachelor of Technology, Computer Science

Amity University, Noida

May 2020

GPA: 3.5/4.0

## SKILLS

**Data Engineering:** Python, SQL, ETL/ELT Pipelines, Data Modeling, Schema Design, Data Quality

**Big Data & Streaming:** Apache Spark, Apache Kafka, Batch & Streaming Processing

**Orchestration & DevOps:** Apache Airflow, dbt, Docker, Git, CI/CD, Terraform

**Cloud & Data Platforms:** AWS (S3, Glue, Redshift), GCP (BigQuery), Microsoft Azure, Snowflake

**Databases & Analytics:** PostgreSQL, MySQL, MongoDB, Tableau, Power BI

## WORK EXPERIENCE

### Summer Intern, *Pixis.ai*, New York

May 2025 – Aug 2025

- Architected batch **ETL** pipelines in **Python** to ingest lead data from external **REST APIs**, standardize schemas, and load data into **Redshift**, reducing manual data processing by 75%.
- Built dimensional **data models** and optimized **SQL** query patterns in Redshift, improving analytical query performance by 40% for marketing analytics workloads.
- Implemented an automated **data quality** framework with null validation, format checks, deduplication, and anomaly detection, improving data accuracy by 20% and eliminating 15 hours/week of manual reconciliation.

### Software Engineer, *Trewinia Infotech*

Nov 2022 – Jul 2024

- Engineered production **ETL** pipelines using **AWS Glue** and **Lambda** to ingest operational data into **PostgreSQL**, processing **1M+ records daily** with automated retries and error handling.
- Designed star schema warehouse tables and tuned **SQL** queries and indexes, reducing report generation time by **50%** for recurring analytical workloads.
- Orchestrated **Python**-based transformation workflows to aggregate, clean, and model raw data into analysis-ready tables powering **15+ dashboards** in **Power BI** and **Tableau**.

### Systems Engineer, *Tata Consultancy Services Ltd*

Feb 2021 – Oct 2022

- Productionized **data processing** services in **Java** and **SQL** supporting **50K+ monthly transactions**, implementing efficient data access patterns and ensuring high service reliability.
- Optimized **database schemas**, **SQL** queries, and indexing strategies to improve data retrieval performance by **35%** for high-traffic application endpoints.
- Automated **CI/CD** pipelines using **Git**, **Jenkins**, and **Python** for build, test, and deployment workflows, reducing release cycle time by **10 hours/week**.

## PROJECT EXPERIENCE

### Digital Audio Archive Processing System, *New York University & New York Public Library*

Dec 2025

- Built a multi-stage **ETL** pipeline using **Python** and **Apache Spark** to ingest and process large-scale digital audio archives, orchestrated with **Apache Airflow** and persisted in **MongoDB** with intermediate staging for audit and recovery.
- Containerized the pipeline using **Docker** and provisioned infrastructure with **Terraform**, achieving a **5× speedup** while processing **500+ hours** of audio data and reducing manual effort by **95%**.

### Startup Analytics Data Pipeline, *NYU Stern School of Business*

May 2025

- Designed **SQL** and **Python**-based transformation workflows to normalize schemas, handle missing data, and prepare **54K+ records** for downstream analytics, staging curated datasets in **Snowflake**.
- Streamlined event-driven ingestion using **Apache Kafka** and built modular transformations with **dbt**, enforcing **data quality** checks and schema consistency for analytical consumers.

## ACHIEVEMENTS

- Wasserman Center Internship Grant (Violet Internship & Research Award)**, *NYU*

Aug 2025

- Microsoft Certified: Azure Fundamentals (AZ-900)**, *Microsoft*

Aug 2022